

Part 1

Loading the required package into Rstudio

```
# First 6 rows of our data set
```

```
head(global_economy)
```

```
## # A tsibble: 6 x 9 [1Y]
```

```
## # Key:   Country [1]
```

```
## Country   Code Year   GDP Growth  CPI Imports Exports Population
```

```
## <fct>     <fct> <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
```

```
## 1 Afghanistan AFG  1960 5377777811.  NA  NA  7.02  4.13  8996351
```

```
## 2 Afghanistan AFG  1961 5488888896.  NA  NA  8.10  4.45  9166764
```

```
## 3 Afghanistan AFG  1962 5466666678.  NA  NA  9.35  4.88  9345868
```

```
## 4 Afghanistan AFG  1963 7511111191.  NA  NA  16.9  9.17  9533954
```

```
## 5 Afghanistan AFG  1964 8000000044.  NA  NA  18.1  8.89  9731361
```

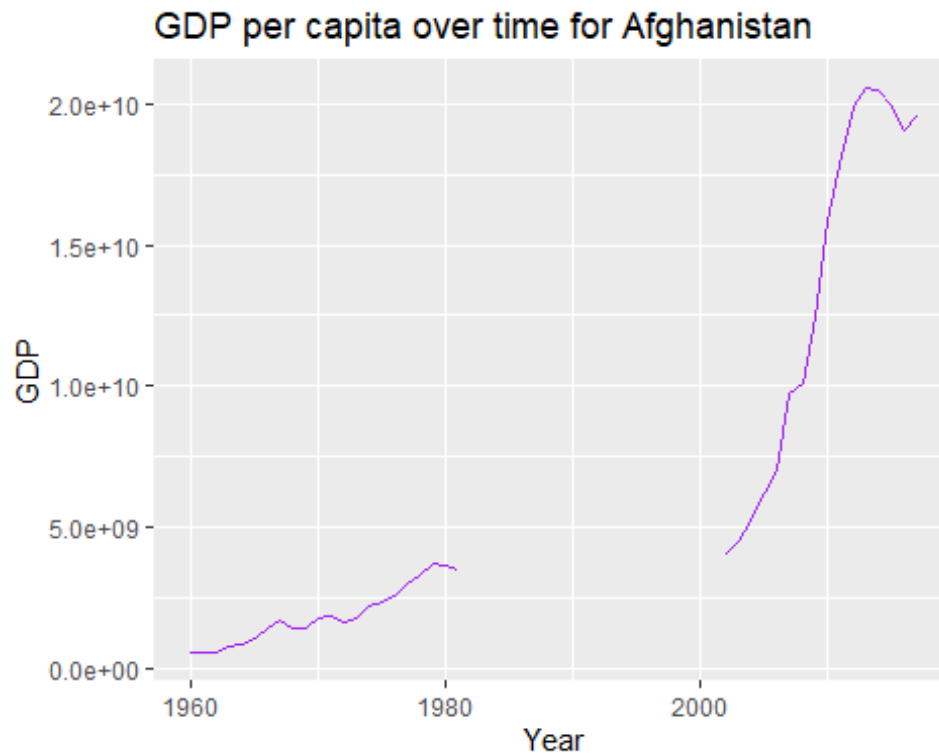
```
## 6 Afghanistan AFG  1965 10066666638.  NA  NA  21.4  11.3  9938414
```

```
# Selecting Vietnam, Uruguay, and Afghanistan countries.
```

```
# Afghanistan
```

```
Afghanistan_data <- global_economy %>% filter(Country=="Afghanistan")
```

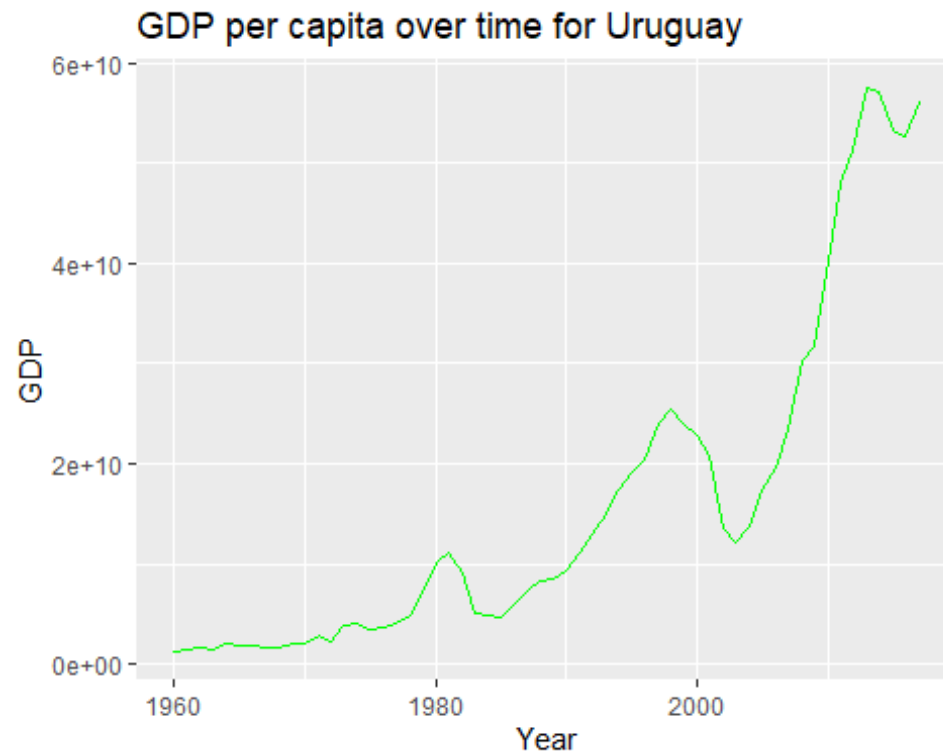
```
ggplot(data = Afghanistan_data,aes(x=Year,y=GDP))+geom_line(col="purple")+ggtitle("GDP  
per capita over time for Afghanistan")
```



#Uruguay

```
Uruguay_data<- global_economy %>% filter(Country=="Uruguay")
```

```
ggplot(data =Uruguay_data,aes(x=Year,y=GDP))+geom_line(col="green")+ggtitle("GDP per  
capita over time for Uruguay")
```

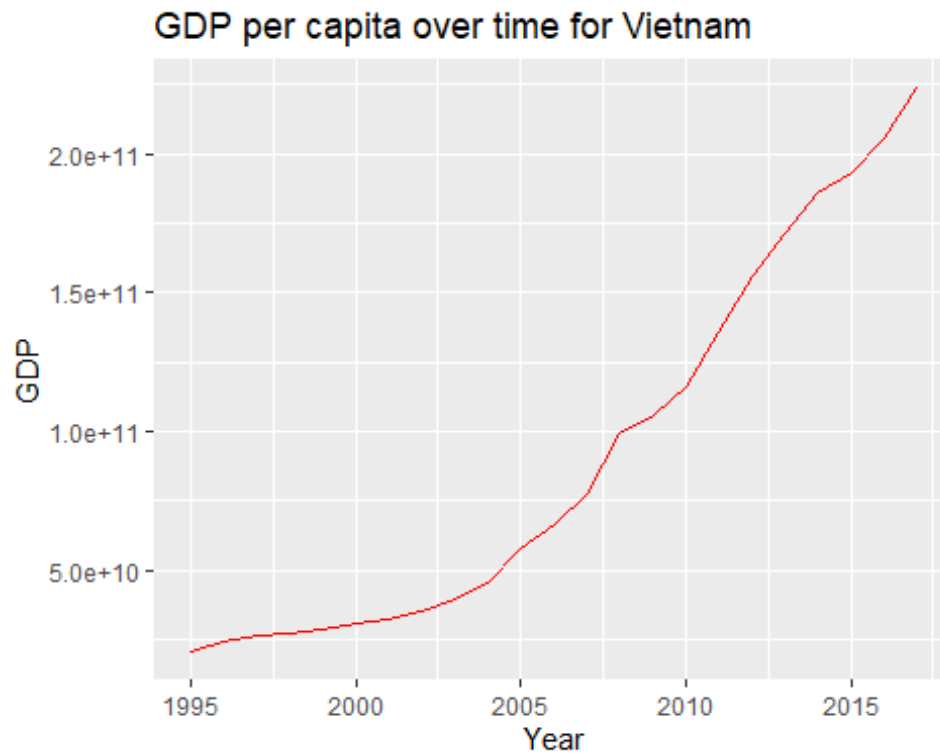


```
#Uruguay
```

```
# Vietnam
```

```
Vietnam_data<-global_economy %>% filter(Country=="Vietnam")
```

```
ggplot(data = na.omit(Vietnam_data),aes(x=Year,y=GDP))+geom_line(col="red")+ggtitle("GDP  
per capita over time for Vietnam")
```



Which has highest GDP per capita?

```
options(scipen = 999)
```

```
sum(Afghanistan_data$GDP,na.rm=T)
```

```
## [1] 253490832378
```

```
sum(Uruguay_data$GDP,na.rm=T)
```

```
## [1] 931336081676
```

```
sum(Vietnam_data$GDP,na.rm=T)
```

```
## [1] 2269838130133
```

```
# Based on the above results, Vietnam has the highest GDP.
```

Checking on the above plots, it is visible that GDP has increased over the years.

PART 2.

Use the canadian_gas data (monthly Canadian gas production in billions of cubic metres, January 1960 – February 2005).

```
str(canadian_gas)


## tbl_ts [542 x 2] (S3: tbl_ts/tbl_df/tbl/data.frame)
##  $ Month : mth [1:542] 1960 Jan, 1960 Feb, 1960 Mar, 1960 Apr, 1960 May, 1960 Jun...
##  $ Volume: num [1:542] 1.43 1.31 1.4 1.17 1.12 ...
## - attr(*, "key")= tibble [1 x 1] (S3: tbl_df/tbl/data.frame)
##  ..$ .rows: list<int> [1:1]
##  .. ..$ : int [1:542] 1 2 3 4 5 6 7 8 9 10 ...
##  .. ..@ ptype: int(0)
## - attr(*, "index")= chr "Month"
##  ..- attr(*, "ordered")= logi TRUE
## - attr(*, "index2")= chr "Month"
## - attr(*, "interval")= interval [1:1] 1M
##  ..@ .regular: logi TRUE

# Canadian_gas data contains 542 observations of 2 variables.

# Plot Volume using autoplot, gg_subseries, gg_season to look at the effect of changing
seasonality over time


autoplot(canadian_gas)+labs(title = "Autoplot of data")
```

```
## Plot variable not specified, automatically selected `.vars = Volume`
```




```
gg_subseries(canadian_gas)+ labs(title = "subseries of data")
```

```
## Plot variable not specified, automatically selected `y = Volume`
```



```
gg_season(canadian_gas) + labs(title = "gg_season of data")
```

```
## Plot variable not specified, automatically selected `y = Volume`
```




Do an STL

decomposition of the data. You will need to choose a seasonal window to allow for the changing shape of the seasonal component.

```
decomposition<-stl(canadian_gas,s.window = "periodic")
```

How does seasonal shape change over time? [Hint: Try plotting the seasonal component using `gg_season().`]

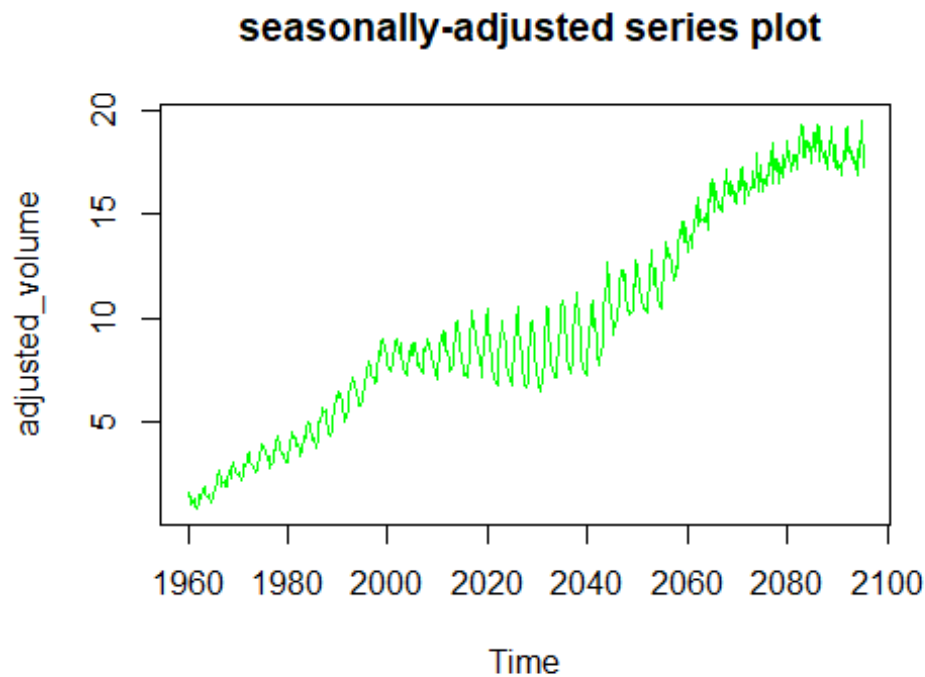
```
plot(decomposition,col="green")
```

Under the seasonal plot above, we can say that the data exhibits a seasonal pattern

Can you produce a plausible seasonally adjusted series?

```
sas<-as.ts(canadian_gas,frequency = 4,start=1990,end=2025)
dec_adjust<-decompose(sas,type="additive")
adjusted_volume<-sas-dec_adjust$seasonal
plot(adjusted_volume,col="green",main=" seasonally-adjusted series plot")
```



#plausible seasonally adjusted series analyses trends by removing all the noises present.

PART 3

Aus Retail Time Series We will use aus_rail dataset Using the code below, get a series (it gets a series randomly by using sample() function):

```
set.seed(1234567)

myseries <- aus_retail %>%
  filter(`Series ID` == sample(aus_retail$`Series ID`,1))
head(myseries)

## # A tsibble: 6 x 5 [1M]
## # Key:   State, Industry [1]
##   State   Industry                Serie~1  Month Turno~2
```

```
## <chr> <chr> <chr> <nth> <dbl>

## 1 Victoria Cafes, restaurants and takeaway food servic~ A33494~ 1982 Apr 85.1
## 2 Victoria Cafes, restaurants and takeaway food servic~ A33494~ 1982 May 85.1
## 3 Victoria Cafes, restaurants and takeaway food servic~ A33494~ 1982 Jun 82.8
## 4 Victoria Cafes, restaurants and takeaway food servic~ A33494~ 1982 Jul 82.1
## 5 Victoria Cafes, restaurants and takeaway food servic~ A33494~ 1982 Aug 81.8
## 6 Victoria Cafes, restaurants and takeaway food servic~ A33494~ 1982 Sep 84.6

## # ... with abbreviated variable names 1: `Series ID`, 2: Turnover

# remover NA's in the series with below:

myseries = myseries %>% filter(!is.na(`Series ID`))

nrow(myseries)

## [1] 441

# rename the column name `Series ID` with MyRandomSeries

rename(myseries, MyRandomSeries = `Series ID`)

## # A tibble: 441 x 5 [1M]

## # Key: State, Industry [1]

## State Industry MyRan~1 Month Turno~2

## <chr> <chr> <chr> <nth> <dbl>

## 1 Victoria Cafes, restaurants and takeaway food servi~ A33494~ 1982 Apr 85.1
## 2 Victoria Cafes, restaurants and takeaway food servi~ A33494~ 1982 May 85.1
## 3 Victoria Cafes, restaurants and takeaway food servi~ A33494~ 1982 Jun 82.8
## 4 Victoria Cafes, restaurants and takeaway food servi~ A33494~ 1982 Jul 82.1
```

```
## 5 Victoria Cafes, restaurants and takeaway food servi~ A33494~ 1982 Aug 81.8
## 6 Victoria Cafes, restaurants and takeaway food servi~ A33494~ 1982 Sep 84.6
## 7 Victoria Cafes, restaurants and takeaway food servi~ A33494~ 1982 Oct 91.7
## 8 Victoria Cafes, restaurants and takeaway food servi~ A33494~ 1982 Nov 97.7
## 9 Victoria Cafes, restaurants and takeaway food servi~ A33494~ 1982 Dec 109.
## 10 Victoria Cafes, restaurants and takeaway food servi~ A33494~ 1983 Jan 94.6
## # ... with 431 more rows, and abbreviated variable names 1: MyRandomSeries,
## # 2: Turnover
```

a) Run a linear regression of Turnover on its trend. Hint: use TSLM() and trend() functions)

```
fit = myseries %>% model(TSLM(Turnover~ trend()))
report(fit)

## Series: Turnover
## Model: TSLM
##
## Residuals:
##   Min     1Q  Median     3Q    Max
## -125.471 -50.951  -9.889  48.598 242.364
##
## Coefficients:
##           Estimate Std. Error t value      Pr(>|t|)
## (Intercept) -23.529     6.121  -3.844    0.000139 ***
## trend()      1.921     0.024  80.057 < 0.0000000000000002 ***
## ---
```

```
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

##

## Residual standard error: 64.17 on 439 degrees of freedom

## Multiple R-squared: 0.9359, Adjusted R-squared: 0.9357

## F-statistic: 6409 on 1 and 439 DF, p-value: < 0.000000000000000222

# checking o the value of Multiple R-squared, it means that our model is 93.59 accurate
```

b) Forecast for next 3 years. What are the values for the next 3 years? Monthly values?

```
x<-myseries$Turnover %>% forecast(h=36)

x

##   Point Forecast   Lo 80   Hi 80   Lo 95   Hi 95
## 442    1014.052 926.9777 1101.126 880.8834 1147.221
## 443    1015.375 918.7325 1112.017 867.5733 1163.176
## 444    1016.697 911.3192 1122.076 855.5354 1177.859
## 445    1018.020 904.5455 1131.495 844.4756 1191.565
## 446    1019.343 898.2828 1140.403 834.1974 1204.488
## 447    1020.666 892.4403 1148.891 824.5620 1216.769
## 448    1021.988 886.9512 1157.025 815.4669 1228.510
## 449    1023.311 881.7643 1164.858 806.8341 1239.788
## 450    1024.634 876.8397 1172.428 798.6022 1250.665
## 451    1025.956 872.1452 1179.768 790.7225 1261.190
## 452    1027.279 867.6549 1186.903 783.1550 1271.403
## 453    1028.602 863.3472 1193.857 775.8667 1281.337
```

## 454	1029.925 859.2039 1200.645 768.8298 1291.019
## 455	1031.247 855.2096 1207.285 762.0209 1300.474
## 456	1032.570 851.3513 1213.789 755.4199 1309.720
## 457	1033.893 847.6175 1220.168 749.0093 1318.776
## 458	1035.215 843.9982 1226.433 742.7739 1327.657
## 459	1036.538 840.4849 1232.591 736.7006 1336.376
## 460	1037.861 837.0698 1238.652 730.7774 1344.944
## 461	1039.184 833.7460 1244.621 724.9939 1353.373
## 462	1040.506 830.5075 1250.505 719.3408 1361.672
## 463	1041.829 827.3488 1256.309 713.8098 1369.848
## 464	1043.152 824.2649 1262.039 708.3932 1377.910
## 465	1044.474 821.2514 1267.698 703.0842 1385.865
## 466	1045.797 818.3042 1273.290 697.8767 1393.718
## 467	1047.120 815.4196 1278.820 692.7649 1401.475
## 468	1048.443 812.5943 1284.291 687.7437 1409.142
## 469	1049.765 809.8251 1289.706 682.8084 1416.722
## 470	1051.088 807.1091 1295.067 677.9545 1424.222
## 471	1052.411 804.4438 1300.378 673.1780 1431.644
## 472	1053.734 801.8267 1305.640 668.4752 1438.992
## 473	1055.056 799.2555 1310.857 663.8427 1446.270
## 474	1056.379 796.7281 1316.030 659.2773 1453.481
## 475	1057.702 794.2426 1321.161 654.7758 1460.628

```
## 476    1059.024 791.7972 1326.252 650.3357 1467.713
```

```
## 477    1060.347 789.3901 1331.304 645.9542 1474.740
```

In the table above Print Forecast column shows the forecasted values of the next 3 years


c) Plot the forecast values with original data

```
ggplot(data =  
myseries,aes(Turnover,x$fitted))+geom_point(col="red")+geom_smooth(col="green")+ggtitle("f  
orecast values Vs original data  
")
```

```
## Don't know how to automatically pick scale for object of type <ts>. Defaulting
```

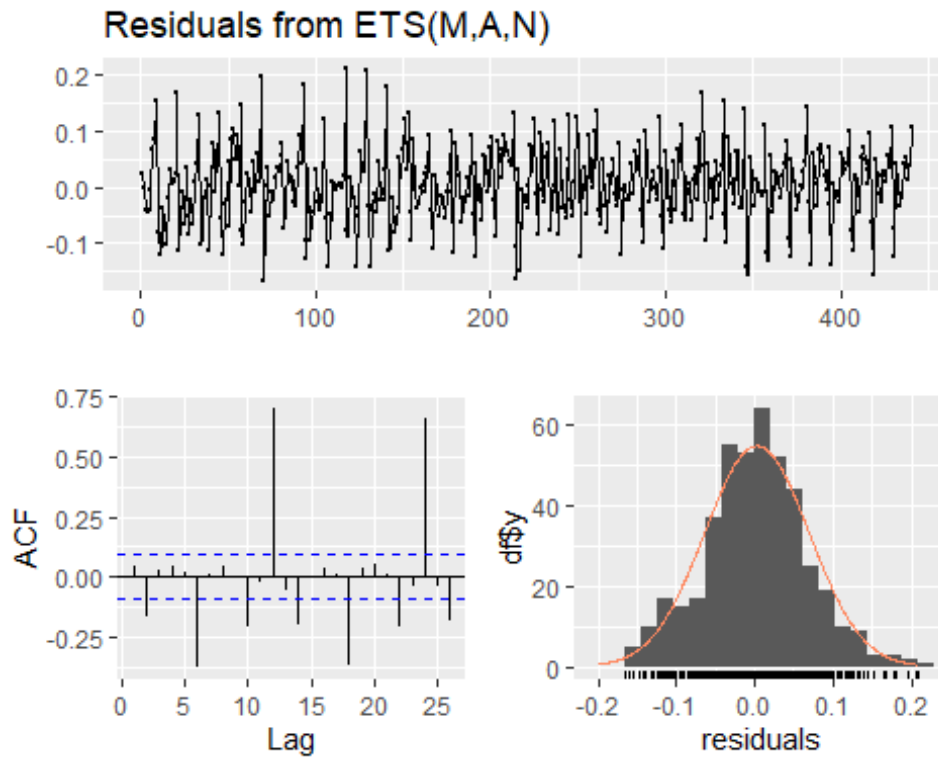
```
## to continuous.
```

```
## `geom_smooth()` using method = 'loess' and formula = 'y ~ x'
```



- d) Get the residuals, does it satisfy requirements for white noise error terms. Hint: `augment()` and `gg_tsresiduals()` functions)

```
checkresiduals(x, plot = T)
```



```
##  
## Ljung-Box test  
##  
## data: Residuals from ETS(M,A,N)  
## Q* = 97.597, df = 6, p-value < 0.000000000000000022  
##  
## Model df: 4. Total lags used: 10
```


By checking the p-value and the plot above, the residuals satisfy the requirements for white noise error terms.